

Each layer consists of attention heads and an MLP

The output of each layer is added to the residual stream, updating the representation

The final output of the model are the probabilities calculated based on these vectors

